

Data Wrangling I

Perform the following operations using Python on any open source dataset (e.g., data.csv) the data types (i.e., character, numeric, integer, factor, and logical) of the variables in the

1. Import all the required Python Libraries.
2. Locate an open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site).
3. Load the Dataset into pandas dataframe.
4. Data Preprocessing: check for missing values in the data using pandas isnull(), describe() function to get some initial statistics. Provide variable descriptions. Types of variables etc. Check the dimensions of the data frame.
5. Data Formatting and Data Normalization: Summarize the types of variables by checking

data set. If variables are not in the correct data type, apply proper type conversions. 6. Turn categorical variables into quantitative variables in Python. In addition to the codes and outputs, explain every operation that you do in the above steps and explain everything that you do to import/read/scrape the data set.

1.1 import all the required Pyhton Libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

Locate an open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site).

```
In [2]: url_link="/home/mca01/Downloads/autodata.csv"
df = pd.read_csv(url_link)
```

```
In [3]: df.head(10)
```

Out[3]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	v
0	0	3	122	alfa-romero	std	two	convertible	
1	1	3	122	alfa-romero	std	two	convertible	
2	2	1	122	alfa-romero	std	two	hatchback	
3	3	2	164	audi	std	four	sedan	
4	4	2	164	audi	std	four	sedan	
5	5	2	122	audi	std	two	sedan	
6	6	1	158	audi	std	four	sedan	
7	7	1	122	audi	std	four	wagon	
8	8	1	158	audi	turbo	four	sedan	
9	9	2	192	bmw	std	two	sedan	

10 rows × 30 columns

In [4]: `df.tail()`

Out[4]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	dri whe
196	196	-1	95	volvo	std	four	sedan	
197	197	-1	95	volvo	turbo	four	sedan	
198	198	-1	95	volvo	std	four	sedan	
199	199	-1	95	volvo	turbo	four	sedan	
200	200	-1	95	volvo	turbo	four	sedan	

5 rows × 30 columns

In [5]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 201 entries, 0 to 200
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            201 non-null   int64
1   symboling              201 non-null   int64
2   normalized-losses      201 non-null   int64
3   make                  201 non-null   object
4   aspiration             201 non-null   object
5   num-of-doors           201 non-null   object
6   body-style             201 non-null   object
7   drive-wheels           201 non-null   object
8   engine-location        201 non-null   object
9   wheel-base             201 non-null   float64
10  length                 201 non-null   float64
11  width                  201 non-null   float64
12  height                 201 non-null   float64
13  curb-weight            201 non-null   int64
14  engine-type            201 non-null   object
15  num-of-cylinders       201 non-null   object
16  engine-size            201 non-null   int64
17  fuel-system            201 non-null   object
18  bore                   201 non-null   float64
19  stroke                 197 non-null   float64
20  compression-ratio      201 non-null   float64
21  horsepower             199 non-null   float64
22  peak-rpm               199 non-null   float64
23  city-mpg               201 non-null   int64
24  highway-mpg            201 non-null   int64
25  price                  201 non-null   float64
26  city-L/100km           201 non-null   float64
27  horsepower-binned      199 non-null   object
28  diesel                 201 non-null   int64
29  gas                    201 non-null   int64
dtypes: float64(11), int64(9), object(10)
memory usage: 47.2+ KB

```

In [6]: `df.describe()`

Out[6]:

	Unnamed: 0	symboling	normalized-losses	wheel-base	length	width
count	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000
mean	100.000000	0.840796	122.000000	98.797015	0.837102	0.91512
std	58.167861	1.254802	31.99625	6.066366	0.059213	0.02918
min	0.000000	-2.000000	65.00000	86.600000	0.678039	0.83750
25%	50.000000	0.000000	101.00000	94.500000	0.801538	0.89027
50%	100.000000	1.000000	122.00000	97.000000	0.832292	0.90972
75%	150.000000	2.000000	137.00000	102.400000	0.881788	0.92500
max	200.000000	3.000000	256.00000	120.900000	1.000000	1.00000

```
In [7]: df.isnull()
```

Out[7]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
196	False	False	False	False	False	False	False	False
197	False	False	False	False	False	False	False	False
198	False	False	False	False	False	False	False	False
199	False	False	False	False	False	False	False	False
200	False	False	False	False	False	False	False	False

201 rows × 30 columns

```
In [8]: df.isnull().sum()
```

```
Out[8]: Unnamed: 0      0
        symboling      0
        normalized-losses 0
        make           0
        aspiration      0
        num-of-doors    0
        body-style      0
        drive-wheels    0
        engine-location 0
        wheel-base      0
        length          0
        width           0
        height          0
        curb-weight     0
        engine-type      0
        num-of-cylinders 0
        engine-size     0
        fuel-system     0
        bore            0
        stroke          4
        compression-ratio 0
        horsepower      2
        peak-rpm        2
        city-mpg        0
        highway-mpg     0
        price           0
        city-L/100km    0
        horsepower-binned 2
        diesel          0
        gas             0
        dtype: int64
```

```
In [9]: df.notnull()
```

Out[9]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	dri whe
0	True	True	True	True	True	True	True	1
1	True	True	True	True	True	True	True	1
2	True	True	True	True	True	True	True	1
3	True	True	True	True	True	True	True	1
4	True	True	True	True	True	True	True	1
...
196	True	True	True	True	True	True	True	1
197	True	True	True	True	True	True	True	1
198	True	True	True	True	True	True	True	1
199	True	True	True	True	True	True	True	1
200	True	True	True	True	True	True	True	1

201 rows × 30 columns

```
In [10]: df.notnull().sum()
```

```
Out[10]: Unnamed: 0      201
        symboling      201
        normalized-losses 201
        make           201
        aspiration      201
        num-of-doors     201
        body-style       201
        drive-wheels     201
        engine-location  201
        wheel-base       201
        length           201
        width            201
        height           201
        curb-weight      201
        engine-type       201
        num-of-cylinders  201
        engine-size      201
        fuel-system       201
        bore             201
        stroke           197
        compression-ratio 201
        horsepower       199
        peak-rpm         199
        city-mpg         201
        highway-mpg      201
        price            201
        city-L/100km     201
        horsepower-binned 199
        diesel           201
        gas              201
        dtype: int64
```

```
In [11]: #calculate the mean value for "stroke" column
avg_stroke = df["stroke"].astype("float").mean(axis = 0)
print("Average of stroke :",avg_stroke)

#replace NaN by mean value in "stroke" column
df["stroke"].replace(np.nan, avg_stroke,inplace = True)
```

Average of stroke : 3.2569035532994857

Calculate the mean value for the 'horsepower' column :

```
In [12]: avg_hp=df["horsepower"].astype("float").mean(axis = 0)
print("Average of stroke :",avg_hp)
```

Average of stroke : 103.39698492462311

```
In [13]: df['horsepower'].replace(np.nan,avg_hp,inplace = True)
```

```
In [14]: from contextlib import nullcontext
df['num-of-doors'].value_counts()
```

```
Out[14]: four      115
        two       86
        num-of-doors, dtype: int64
```

```
In [15]: df['num-of-doors'].value_counts().idxmax()
```

```
Out[15]: 'four'
```

```
In [16]: # replace the missing 'num-of-door' values by most frequent  
df['num-of-doors'].replace(np.nan, "four" , inplace=True)  
  
#simply drop whole row with nan in "Horsepower-binned" column  
df.dropna(subset=['horsepower-binned'], axis=0 , inplace=True)  
  
#reset index, because we dropped two rows  
df.reset_index(drop=True, inplace=True)
```

```
In [17]: df.isnull().sum()
```

```
Out[17]: Unnamed: 0      0  
symboling      0  
normalized-losses  0  
make          0  
aspiration     0  
num-of-doors   0  
body-style     0  
drive-wheels   0  
engine-location 0  
wheel-base    0  
length        0  
width         0  
height        0  
curb-weight    0  
engine-type    0  
num-of-cylinders 0  
engine-size    0  
fuel-system    0  
bore          0  
stroke        0  
compression-ratio 0  
horsepower     0  
peak-rpm      0  
city-mpg       0  
highway-mpg    0  
price         0  
city-L/100km   0  
horsepower-binned 0  
diesel        0  
gas          0  
dtype: int64
```

DATA STANDARDIZATION : It is process of transforming data into common format which allows the researcher to make meaningful comparison

```
In [18]: df['city-L/100km']=235/df['city-mpg']  
df.head()
```


Out[18]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	v
0	0	3	122	alfa-romero	std	two	convertible	
1	1	3	122	alfa-romero	std	two	convertible	
2	2	1	122	alfa-romero	std	two	hatchback	
3	3	2	164	audi	std	four	sedan	
4	4	2	164	audi	std	four	sedan	

5 rows × 30 columns

```
In [19]: df['highway-L/100km']=235/df["highway-mpg"]  
df.head()
```

Out[19]:

	Unnamed: 0	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	v
0	0	3	122	alfa-romero	std	two	convertible	
1	1	3	122	alfa-romero	std	two	convertible	
2	2	1	122	alfa-romero	std	two	hatchback	
3	3	2	164	audi	std	four	sedan	
4	4	2	164	audi	std	four	sedan	

5 rows × 31 columns

DATA NORMALIZATION : It is process of transforming several values into similar range

```
In [20]: df['length']=df['length']/df['length'].max()  
df['width']=df['width']/df['width'].max()  
df['height']=df['height']/df['height'].max()
```

```
In [21]: df[['length','width','height']].head()
```

Out[21]:

	length	width	height
--	--------	-------	--------

0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

INDIACTOR VARIABLE : Indicator variable or dummy variable are used to label numerical variable used to label categories

In [22]: `df.columns`

Out[22]: Index(['Unnamed: 0', 'symboling', 'normalized-losses', 'make', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price', 'city-L/100km', 'horsepower-binned', 'displacement', 'gas', 'highway-L/100km'], dtype='object')

In [23]: `df['aspiration'].value_counts()`

Out[23]:

std	163
turbo	36

Name: aspiration, dtype: int64

In [24]: `dummy_var_1=pd.get_dummies(df['aspiration'])`
`dummy_var_1.head()`

Out[24]:

	std	turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

In [25]: `df=pd.concat([df,dummy_var_1], axis=1)`
`df.drop('aspiration',axis = 1 , inplace = True)`

In [26]: `df.head()`

Out[26]:

	Unnamed: 0	symboling	normalized-losses	make	num-of-doors	body-style	drive-wheels	eng loca
0	0	3	122	alfa-romero	two	convertible	rwd	1
1	1	3	122	alfa-romero	two	convertible	rwd	1
2	2	1	122	alfa-romero	two	hatchback	rwd	1
3	3	2	164	audi	four	sedan	fwd	1
4	4	2	164	audi	four	sedan	4wd	1

5 rows × 32 columns

The last columns are indicator variable which are represented by 0's and 1's

In [27]: `df.columns`

Out[27]: Index(['Unnamed: 0', 'symboling', 'normalized-losses', 'make', 'num-of-door
s',
 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'leng
th',
 'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',
 'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',
 'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price',
 'city-L/100km', 'horsepower-binned', 'diesel', 'gas', 'highway-L/100
km',
 'std', 'turbo'],
 dtype='object')


BINNING: It is process of transforming continous data into discrete categorical
'bins' for group analysis

In [28]: `df ["horsepower"]=df ["horsepower"].astype(float, copy=True)`

In [29]: `%matplotlib inline
import matplotlib.pyplot as plt
from matplotlib import pyplot
import numpy as np

plt.matplotlib.pyplot.hist(df['wheel-base'])
plt.matplotlib.pyplot.xlabel('wheel-base')
plt.matplotlib.pyplot.ylabel('count')
plt.matplotlib.pyplot.title('wheel-base bins')`

Out[29]: Text(0.5, 1.0, 'wheel-base bins')



```
In [30]: bins = np.linspace(min(df['wheel-base']),max(df['wheel-base']),4)
bins
```

```
Out[30]: array([ 86.6          ,  98.03333333, 109.46666667, 120.9          ])
```

```
In [ ]:
```